

Phase 1: Standardizing SSD Endurance Metrics in smartmontools

Introduction

The `smartmontools` software is a critical tool for monitoring the health of storage devices, including Solid State Drives (SSDs). However, SSD vendors use inconsistent attribute names for endurance metrics, such as "DriveLife_Remaining%", "Percent_Life_Used", and "Wear_Leveling_Count". This variability complicates comparing drive health across different devices, especially in large fleets or when integrating with tools like Prometheus exporters.

This project, part of Phase 1 for the Google Summer of Code (GSOC) SMART project, aims to enhance `smartmontools` by standardizing the reporting of SSD endurance metrics. By improving the detection and normalization of these attributes in the `ataprint.cpp` file, i introduce a unified "endurance_used" metric in the JSON output, making it easier for users to monitor SSD health consistently.

Proposed Solution

To address the inconsistency in SSD endurance reporting, the following enhancements were implemented:

1. **Expanded Regular Expression (Regex):** The `endurance_regex` was updated to include a broader range of attribute names found in the `drivedb.h` file, ensuring comprehensive coverage of vendor-specific endurance metrics.
2. **Normalization Logic:** A logic was introduced to normalize attribute values to a "percentage used" format. Attributes indicating "remaining" life are inverted ($100 - \text{value}$), while "used" life attributes are reported directly.
3. **Secondary Regex for Flexibility:** A `remaining_regex` was added to dynamically identify attributes representing "remaining" life, improving maintainability and reducing hardcoded checks.

Implementation Details

The changes were implemented in the `ataprint.cpp` file of the `smartmontools` repository. Below are the key components of the implementation.

Regular Expression Definitions

Two regular expressions were defined to handle endurance-related attributes:

```
static const regular_expression endurance_regex(
    "SSD_Life_Left.*|Wear_Leveling.*|Media_Wearout_Indicator.*|DriveLife_Used%.*|DriveLife_Re
    maining%.*|Drive_Life_Used.*|Drive_Life_Remaining.*|
    "Lifetime_.*|Percent_Life_Used.*|Percent_Life_Remaining.*|Remaining_Life.*|End_of_Life.*"
);

static const regular_expression remaining_regex(
    ".*Left.*|.Remaining.*|.Media_Wearout_Indicator.*|.End_of_Life.*"
);
```

- **endurance_regex**: Matches attribute names related to SSD endurance, such as "SSD_Life_Left", "Wear_Leveling_Count", and "Percent_Life_Used".
- **remaining_regex**: Identifies attributes that indicate "remaining" life, such as those containing "Remaining", "Left", or "End_of_Life".

Normalization Logic

The logic checks if an attribute's name matches **endurance_regex**. If it does, it further checks **remaining_regex** to determine if the value represents "remaining" life, requiring inversion.

```
if (id >= 100 && endurance_regex.full_match(name)) {
    if (remaining_regex.full_match(name)) {
        jglb["endurance_used"]["current_percent"] = (normval <= 100 ? 100 - normval : 0);
    } else {
        jglb["endurance_used"]["current_percent"] = normval;
    }
    return;
}
```

- If the attribute matches **remaining_regex**, the value is inverted (e.g., 99% remaining becomes 1% used).
- Otherwise, the raw value (**normval**) is used directly as the percentage used.
- The result is stored in the JSON output under **jglb["endurance_used"]["current_percent"]**.

Attribute Classification

The following table lists endurance-related attributes identified in `drivedb.h` and their classification as "used" or "remaining":

Attribute Name	Classification	Notes / Reasoning
DriveLife_Used%	Used	Reports percent of SSD life used directly
Drive_Life_Used%	Used	Same as above, vendor variation
Percent_Life_Used	Used	Explicitly denotes percentage used
Percent_Lifetime_Used	Used	Common vendor label for usage
Perc_Rated_Life_Used	Used	Rated endurance used, directly reportable
Wear_Leveling_Count	Used	Often increases with wear, vendor-defined
Media_Wearout_Indicator	Used	Some vendors count down, but normalized via regex
SSD_LifeLeft(0.01%)	Used	Granular % used value
SSD_Life_Left_Perc	Used	Vendor reports this as % used
DriveLife_Remaining%	Remaining	Indicates remaining life, invert to get used
Drive_Life_Remaining%	Remaining	Same as above with different naming
Lifetime_Remaining%	Remaining	Needs inversion (100 - value)
Remaining_Life	Remaining	General label for life left
Remaining_Lifetime_Perc	Remaining	Percent of life left
PCT_Life_Remaining	Remaining	Abbreviated form, invert
Perc_Rated_Life_Remain	Remaining	Manufacturer rating for remaining life
Percent_Life_Remaining	Remaining	Indicates remaining percentage
Percent_Lifetime_Remain	Remaining	Same as above, more specific
Sys_Percent_Life_Remain	Remaining	System-derived remaining estimate
SSD_Life_Left	Remaining	Common vendor label for remaining life
Lifetime_Left	Remaining	Another synonym for remaining life
End_of_Life	Remaining	Indicates near-zero remaining life
Life_Remaining_Percent	Ambiguous	Label unclear — may vary by vendor
Life_Curve_Status	Ambiguous	Possibly vendor-specific health flag
Lifetime_NAND_Prg_GiB	Ambiguous	Raw programmed data, not a normalized value
Lifetime_Nand_Writes	Ambiguous	Cumulative writes — not directly percentage
Lifetime_Writes	Ambiguous	Generic, unit undefined
Lifetime_Write_AmpFctr	Ambiguous	Write amplification — quality, not quantity
Wear_Range_Delta	Ambiguous	Spread of wear — not absolute wear level
Reallocated_Sector_Ct	Ambiguous	May correlate with wear, but not always
Lifetime_Die_Failure_Ct	Not Endurance	Tracks hardware failures
Lifetime_UECC_Ct	Not Endurance	Correctable ECC errors
Lifetime_PS3_Exit_Ct	Not Endurance	Power state tracking
Lifetime_PS4_Entry_Ct	Not Endurance	Power state entry count

Attribute Name	Classification	Notes / Reasoning
Lifetime_Reads_GiB	Not Endurance	Total data read, not normalized
Lifetime_Rds_Frm_Hst_GB	Not Endurance	Same as above, alternate format
Lifetime_Writes_GiB	Not Endurance	Raw cumulative writes
Lifetime_Wts_Frm_Hst_GB	Not Endurance	Host-to-drive writes
Lifetime_Wts_To_Flsh_GB	Not Endurance	Internal flash write total
Lifetime_Nand_Gb	Not Endurance	Capacity-related
Timed_Workld_Media_Wear	Not Endurance	Workload test metric
Workld_Media_Wear_Indic	Not Endurance	Workload test metric
Perc_Write/Erase_Count	Not Endurance	Not an endurance percentage

Frequently Used Attributes (High Priority):

Based on vendor frequency and presence in modern SSDs, the following endurance attributes are commonly observed and should be prioritized:

Attribute	Why It's Important
Media_Wearout_Indicator	Found in most Intel, Micron, and some Samsung drives; counts down from 100 to 0
Wear_Leveling_Count	Present in many SATA SSDs (e.g., Samsung, WD Green); counts erase cycles or lifetime
SSD_Life_Left	Used by Kingston, WD, and Crucial drives to report % remaining; often ID 231
DriveLife_Used %	Samsung PM series; reported as % used directly

`Percent_Lifetime_Remain`

Common in enterprise Seagate, Dell, and some NVMe drives

`End_of_Life`

Used by some vendors as a final countdown flag; may appear when drive is >90% used

These account for the majority of endurance indicators across Samsung, Kingston, and Seagate enterprise SSDs (based on sampling from `drivedb.h` and observed test cases).

Testing and Validation on real ssd

The changes were tested by rebuilding `smartmontools` and running the `smartctl` utility:

1. Build Process:

- Executed `make clean` and `make` to rebuild the software.
- Ensured necessary build files (`autogen.sh`, `configure`, `Makefile.am`) were present.

2. Testing Command:

- Ran `smartctl -a /dev/diskX` to retrieve SMART data and inspect the JSON output.

Expected output not yet got:

Expected output includes a standardized `"endurance_used"` field, e.g.:

```
{
  "endurance_used": {
    "current_percent": 73
  }
}
```

3. Attribute Coverage:

Analyzed `drivedb.h` using the command:

```
egrep -o 'label = "[^"]*"' drivedb.h | awk -F'"' '{print $2}' | sort | uniq
```

This ensured all relevant attributes were included in `endurance_regex`.

In case of leaving out some other attribute label other than the ones specified i used this command for this purpose: (maybe a better approach in order to not leave out any attribute)

```
egrep -i 'Wear|Life|Lifetime|Remain|Media_Wearout|Percent' drivedb.h | \
grep -v '^//' | \
awk -F',' '{print $3}' | \
grep -E '0x[0-9A-Fa-f]+' | \
sort | uniq
```

`egrep -i 'Wear|Life|Lifetime|Remain|Media_Wearout|Percent'`: broader keyword match

`grep -v '^//'`: ignore commented-out lines

`awk -F',' '{print $3}'`: pick the 3rd comma-separated field (usually SMART ID)

`grep -E '0x[0-9A-Fa-f]+'`: ensure it's a hex SMART ID

`sort | uniq`: deduplicate

But not got any output as such so ran only the `egrep` to check **what lines are matched at all so showed the below keywords**

The output i got after running this command on terminal: (these are more than 100+)

```
(base) anusha@Anushas-MacBook-Air src % egrep -i
'Wear|Life|Lifetime|Remain|Media_Wearout|Percent' drivedb.h
```

```
"-v 177,raw48,Wear_Leveling_Count,SSD "
"-v 233,raw48,Media_Wearout_Indicator,SSD "
"-v 215,raw48,Current_TRIM_Percent "
"-v 161,raw48,Spare_Blocks_Remaining "
"-v 248,raw48,Perc_Rated_Life_Remain "
"-v 249,raw48,Spares_Remaining_Perc "
"-v 231,raw48,Life_Remaining_Percent "
"-v 248,raw48,SSD_Remaining_Life_Perc "
"-v 248,raw48,Remaining_Life "
"-v 231,raw48,Lifetime_Left "
```

```

"-v 231,raw48,Lifetime_Left "
"-v 173,raw48,Wear_Leveling_Count " // ]
"-v 173,raw48,Wear_Leveling_Count " // ]
"-v 231,raw48,Percent_Lifetime_Remain "
"-v 17,raw48,Remaining_Spare_Blocks "
"-v 202,raw48,Percent_Lifetime_Used "
"-v 248,raw48,Percent_Lifetime_Remain "
"-v 249,raw48,Spare_Blocks_Remaining " // same as ID 17 (Remaining_Spare_Blocks)
"-v 17,raw48,Remaining_Spare_Blocks "
"-v 202,raw48,Percent_Lifetime_Used "
"-v 248,raw48,PCT_Life_Remaining "
"-v 249,raw48,Spare_Block_Remaining "
"-v 173,raw48,Wear_Leveling_Count "
"-v 202,raw48,Percent_Lifetime_Used "
"-v 173,raw48,Wear_Leveling_Count "
"-v 202,raw48,Perc_Rated_Life_Used "
"-v 173,raw48,Wear_Leveling_Count "
"-v 202,raw48,Perc_Rated_Life_Used "
"-v 202,raw48,Percent_Lifetime_Remain " // norm = max(100-raw,0); raw =
percent_lifetime_used
"-v 231,raw48,SSD_Life_Left "
"-v 170,raw48,Reserved_Block_Pct " // Percentage of remaining reserved blocks available
//"-v 180,raw48,Unused_Rsvd_Blks_Cnt_Tot " // absolute count of remaining reserved blocks
available
"-v 202,raw48,Percent_Lifetime_Remain " // Remaining endurance, trips at 10%
"-v 213,raw48,Integ_Scan_Progress " // Current is percentage, raw is absolute number of
superblocks scanned by the current integrity scan
"-v 202,raw48,Percent_Lifetime_Remain "
"-v 231,raw48,SSD_Life_Left "
"-v 241,raw48,Lifetime_Writes_GiB "
"-v 242,raw48,Lifetime_Reads_GiB"
"-v 230,raw48,Media_Wearout_Indicator "
"-v 241,raw48,Lifetime_Writes_GiB "
"-v 242,raw48,Lifetime_Reads_GiB"
"-v 177,raw48,Wear_Range_Delta "
"-v 230,raw48,Life_Curve_Status "
"-v 231,raw48,SSD_Life_Left "
"-v 241,raw48,Lifetime_Writes_GiB "
"-v 242,raw48,Lifetime_Reads_GiB"
// 0729 - remaining in block life. In 0828 remaining is normalized to 100% then decreases
"-v 161,raw48,Spares_Remaining "
"-v 169,raw48,Lifetime_Remaining% "
"-v 248,raw48,Lifetime_Remaining% " // later then 0409 FW.
"-v 249,raw48,Spares_Remaining_Perc " // later then 0409 FW.

```

```

"-v 231,raw48,SSD_Life_Left "
"-v 241,raw48,Lifetime_Writes_GiB "
"-v 242,raw48,Lifetime_Reads_GiB "
"-v 231,raw48,SSD_Life_Left "
"-v 241,raw48,Lifetime_Writes_GiB "
"-v 242,raw48,Lifetime_Reads_GiB "
"-v 209,raw64,Remaining_Lifetime_Perc "
"-v 232,raw48,Lifetime_Writes " // LBA?
//"-v 233,raw48,Media_Wearout_Indicator"
"-v 233,raw48,Remaining_Lifetime_Perc "
"-v 233,raw48,Remaining_Lifetime_Perc "
"-v 169,raw48,Remaining_Lifetime_Perc "
//"-v 177,raw48,Wear_Leveling_Count "
"-v 248,raw48,Remaining_Life "
"-v 249,raw48,Spare_Blocks_Remaining"
"-v 169,raw48,Remaining_Lifetime_Perc "
"-v 232,raw48,Spares_Remaining_Perc "
//"-v 233,raw48,Media_Wearout_Indicator"
//"-v 233,raw48,Media_Wearout_Indicator"
"-v 226,raw48,Workld_Media_Wear_Indic " // Timed Workload Media Wear Indicator
(percent*1024)
"-v 227,raw48,Workld_Host_Reads_Perc " // Timed Workload Host Reads Percentage
//"-v 233,raw48,Media_Wearout_Indicator"
"-v 226,raw48,Workld_Media_Wear_Indic "
"-v 226,raw48,Workld_Media_Wear_Indic " // Timed Workload Media Wear Indicator
(percent*1024)
"-v 227,raw48,Workld_Host_Reads_Perc " // Timed Workload Host Reads Percentage
//"-v 233,raw48,Media_Wearout_Indicator "
"-v 226,raw48,Workld_Media_Wear_Indic " // Timed Workload Media Wear Indicator
(percent*1024)
"-v 227,raw48,Workld_Host_Reads_Perc " // Timed Workload Host Reads Percentage
//"-v 233,raw48,Media_Wearout_Indicator "
"-v 226,raw48,Workld_Media_Wear_Indic " // Timed Workload Media Wear Indicator
(percent*1024)
"-v 227,raw48,Workld_Host_Reads_Perc " // Timed Workload Host Reads Percentage
//"-v 233,raw48,Media_Wearout_Indicator "
//"-v 233,raw48,Media_Wearout_Indicator"
"-v 226,raw48,Workld_Media_Wear_Indic "
//"-v 233,raw48,Media_Wearout_Indicator "
"-v 226,raw48,Workld_Media_Wear_Indic "
//"-v 233,raw48,Media_Wearout_Indicator "
"-v 226,raw48,Workld_Media_Wear_Indic "
//"-v 233,raw48,Media_Wearout_Indicator "
//"-v 233,raw48,Media_Wearout_Indicator "

```



```

"-v 226,raw48,Workld_Media_Wear_Indic "
"-v 226,raw48,Workld_Media_Wear_Indic " // Timed Workload Media Wear Indicator
(percent*1024)
"-v 227,raw48,Workld_Host_Reads_Perc " // Timed Workload Host Reads Percentage
//"-v 233,raw48,Media_Wearout_Indicator "
"-v 226,raw48,Workld_Media_Wear_Indic " // Timed Workload Media Wear Indicator
(percent*1024)
"-v 227,raw48,Workld_Host_Reads_Perc " // Timed Workload Host Reads Percentage
//"-v 233,raw48,Media_Wearout_Indicator "
"-v 226,raw48,Workld_Media_Wear_Indic "
"-v 226,raw48,Workld_Media_Wear_Indic "
//"-v 233,raw48,Media_Wearout_Indicator "
"-v 226,raw48,Workld_Media_Wear_Indic "
"-v 226,raw48,Workld_Media_Wear_Indic " // Timed Workload Media Wear Indicator
(percent*1024)
"-v 227,raw48,Workld_Host_Reads_Perc " // Timed Workload Host Reads Percentage
//"-v 233,raw48,Media_Wearout_Indicator "
"-v 202,raw48,End_of_Life "
"-v 226,raw48,Workld_Media_Wear_Indic " // Timed Workload Media Wear Indicator
(percent*1024)
"-v 227,raw48,Workld_Host_Reads_Perc " // Timed Workload Host Reads Percentage
"-v 245,raw48,Percent_Life_Remaining"
"-v 226,raw48,Workld_Media_Wear_Indic " // Timed Workload Media Wear Indicator
(percent*1024)
"-v 227,raw48,Workld_Host_Reads_Perc " // Timed Workload Host Reads Percentage
//"-v 233,raw48,Media_Wearout_Indicator "
"-v 202,raw48,End_of_Life "
"-v 226,raw48,Workld_Media_Wear_Indic " // Timed Workload Media Wear Indicator
(percent*1024)
"-v 227,raw48,Workld_Host_Reads_Perc " // Timed Workload Host Reads Percentage
"-v 245,raw48,Percent_Life_Remaining"
"-v 226,raw48,Workld_Media_Wear_Indic "
"-v 226,raw48,Workld_Media_Wear_Indic "
//"-v 177,raw48,Wear_Leveling_Count "
"-v 231,raw48,SSD_Life_Left "
// (1.04/5 Firmware self-test log lifetime unit is bogus, possibly 1/256 hours)
//"-v 177,raw48,Wear_Leveling_Count "
//"-v 233,raw48,Media_Wearout_Indicator " // MS6/1.03
"-v 173,raw48,Wear_Leveling_Count " // CM871
//"-v 177,raw48,Wear_Leveling_Count "
//"-v 233,raw48,Media_Wearout_Indicator " // PM851, 840
"-v 245,raw48,Timed_Workld_Media_Wear " // PM863, PM893
"-v 201,raw48,Lifetime_Remaining% "
"-v 13,raw48,Lifetime_UECC_Ct "

```

```

"-v 32,raw48,Lifetime_Write_AmpFctr "
"-v 173,raw48,Percent_Life_Used "
"-v 175,raw48,Lifetime_Die_Failure_Ct "
"-v 177,raw48,Lifetime_Remaining% "
"-v 178,raw48,SSD_LifeLeft(0.01%) "
"-v 196,raw48,Lifetime_Retried_Blks_Ct "
"-v 233,raw48,Lifetime_Nand_Writes "
"-v 245,raw48,Drive_Life_Remaining% "
"-v 253,raw48,SPI_Test_Remaining "
"-v 170,raw48,Reserve_Blks_Remaining "
"-v 173,raw48,Drive_Life_Used% "
"-v 177,raw48,DriveLife_Remaining% "
"-v 178,raw48,SSD_Life_Left "
"-v 245,raw48,DriveLife_Used% "
"-v 169,raw48,Remaining_Lifetime_Perc "
//"-v 177,raw48,Wear_Leveling_Count "
"-v 231,raw48,SSD_Life_Left " // KINGSTON SKC600256G/S4500105
"-v 231,raw48,SSD_Life_Left "
"-v 241,raw48,Lifetime_Writes_GiB "
"-v 242,raw48,Lifetime_Reads_GiB "
"-v 169,raw48,Remaining_Lifetime_Perc "
"-v 17,raw48,Spare_Blocks_Remaining " // spec DeclID is wrong, HexID is right
"-v 202,raw48,Perc_Rated_Life_Used "
"-v 231,raw48,Perc_Rated_Life_Remain "
"-v 251,raw48,Min_Spares_Remain_Perc " // percentage of the total number of spare blocks
available
//"-v 177,raw48,Wear_Leveling_Count "
"-v 181,raw48,Sys_Percent_Life_Remain "
"-v 202,raw48,Percent_Lifetime_Remain "
"-v 231,raw48,Percent_Lifetime_Remain "
"-v 233,raw48,Percent_Lifetime_Remain "
"-v 103,raw48,Remaining_Energy_Storg "
"-v 173,raw48,Wear_Leveling_Count "
"-v 231,raw48,Percent_Lifetime_Remain "
"-v 173,raw48,Wear_Leveling_Count "
//"-v 177,raw48,Wear_Leveling_Count "
"-v 201,raw48,Percent_Lifetime_Remain "
"-v 231,raw48,SSD_Life_Left "
"-v 231,raw48,SSD_Life_Left "
"-v 234,raw48,Lifetime_NAND_Prg_GiB " // ?
"-v 235,raw48,Lifetime_Writes_GiB "
"-v 241,raw48,Lifetime_NAND_Prg_GiB "
"-v 242,raw48,Lifetime_Reads_GiB "
"-v 245,raw48,SSD_Life_Left "

```

```

"-v 231,raw48,SSD_Life_Left_Perc "
"-v 234,raw48,Lifetime_Nand_Gb "
"-v 102,raw48,Lifetime_PS4_Entry_Ct "
"-v 103,raw48,Lifetime_PS3_Exit_Ct "
"-v 177,raw16,Wear_Range_Delta "
"-v 230,raw56,Drv_Life_Protect_Status "
"-v 231,hex56,SSD_Life_Left "
"-v 233,raw48,Lifetime_Wts_To_Flsh_GB "
"-v 241,raw48,Lifetime_Wts_Frm_Hst_GB "
"-v 242,raw48,Lifetime_Rds_Frm_Hst_GB "
"-v 17,raw48,Spare_Blocks_Remaining "
"-v 177,raw16,Wear_Range_Delta "
"-v 231,raw48,SSD_Life_Left "
"-v 102,raw48,Lifetime_PS4_Entry_Ct "
"-v 103,raw48,Lifetime_PS3_Exit_Ct "
"-v 177,raw16,Wear_Range_Delta "
"-v 230,raw56,Drv_Life_Protect_Status "
"-v 231,hex56,SSD_Life_Left "
"-v 233,raw48,Lifetime_Wts_To_Flsh_GB "
"-v 241,raw48,Lifetime_Wts_Frm_Hst_GB "
"-v 242,raw48,Lifetime_Rds_Frm_Hst_GB "
"-v 230,hex48,Media_Wearout_Indicator " // Maybe hex16

```

```

// --- Regex to match all known SSD endurance/lifetime-related SMART attributes ---
// Covers vendor variations like 'SSD_Life_Left', 'DriveLife_Used%',
'Percent_Lifetime_Remain', etc.
static const regular_expression endurance_regex(
    "SSD_Life_Left.*|Wear_Leveling.*|"
    "DriveLife_Remaining%|DriveLife_Used%|"
    "Drive_Life_Remaining%|Drive_Life_Used%|"
    "SSD_Life_Left_Perc|SSD_Remaining_Life_Perc|"
    "Percent_Life_Remaining|Percent_Life_Used|"
    "PCT_Life_Remaining|Perc_Rated_Life_Remain|"
    "Perc_Rated_Life_Used|Remaining_Life|"
    "Lifetime_Remaining%|Lifetime_Left|"
    "Media_Wearout_Indicator|Percent_Lifetime_Remain|"
    "Percent_Lifetime_Used|End_of_Life"
);

// --- Regex to detect whether the attribute expresses remaining life ---
// If matched, value is inverted (100 - normval) to give % life used

```

```

static const regular_expression
remaining_regex(".*Remaining.*|.*Left.*|.*Remain.*|.*End_of_Life.*");

// --- Endurance Normalization Logic ---
// If a SMART attribute matches the endurance regex, normalize its value
// Store the result under the JSON key: "endurance_used.current_percent"
if (id >= 100 && endurance_regex.full_match(name)) {
    bool isRemaining = remaining_regex.full_match(name); // Check if it's a 'remaining
life' type
    jglb["endurance_used"]["current_percent"] = isRemaining
        ? (normval <= 100 ? 100 - normval : 0) // Invert value if needed
        : normval; // Else use as-is
    return; // Exit after handling this attribute
}

// --- Temperature Attribute Parsing ---
// Matches all temperature-related labels like "Temperature_Celsius", "Temp_Internal",
etc.
static const regular_expression temperature_regex(".*[Tt]emperature.*|.*[Tt]emp.*");

// If attribute name matches temperature regex, store it under "temperature_celsius"
if (temperature_regex.full_match(name)) {
    jglb["temperature_celsius"] = normval; // Direct assignment as it's already
normalized
    return;
}

// --- LBA Written / NAND Write Attribute Parsing ---
// Matches labels like "Host_Writes", "LBAs_Written", "Total_Writes_GiB", etc.
static const regular_expression lba_written_regex(

".*LBAs.*Written.*|.*Host.*Writes.*|.*Writes.*GiB.*|.*Total.*Writes.*|.*NAND.*Writes.*
|.*Program_Page_Count.*"
);

// If matched, store raw value under "lbas_written"
if (lba_written_regex.full_match(name)) {
    jglb["lbas_written"] = normval; // Unit conversion (e.g., to GiB) can be added
later if needed
    return;
}

```

```
// --- Fallback Logging ---
#ifdef NDEBUG
std::cerr << "Unmatched SMART attribute: " << name << " (ID: " << id << ", normval: "
<< normval << ")" << std::endl;
#endif
```

// --- Regex to match all known SSD endurance/lifetime-related SMART attributes ---
// Covers vendor variations like 'SSD_Life_Left', 'DriveLife_Used%', 'Percent_Lifetime_Remain',
etc.

```
static const regular_expression endurance_regex(
    "SSD_Life_Left.*|Wear_Leveling.*|"
    "DriveLife_Remaining%|DriveLife_Used%|"
    "Drive_Life_Remaining%|Drive_Life_Used%|"
    "SSD_Life_Left_Perc|SSD_Remaining_Life_Perc|"
    "Percent_Life_Remaining|Percent_Life_Used|"
    "PCT_Life_Remaining|Perc_Rated_Life_Remain|"
    "Perc_Rated_Life_Used|Remaining_Life|"
    "Lifetime_Remaining%|Lifetime_Left|"
    "Media_Wearout_Indicator|Percent_Lifetime_Remain|"
    "Percent_Lifetime_Used|End_of_Life"
);
```

// --- Regex to detect whether the attribute expresses remaining life ---
// If matched, value is inverted (100 - normval) to give % life used
static const regular_expression
remaining_regex(".*Remaining.*.*Left.*.*Remain.*.*End_of_Life.*");

// --- Endurance Normalization Logic ---
// If a SMART attribute matches the endurance regex, normalize its value
// Store the result under the JSON key: "endurance_used.current_percent"
if (id >= 100 && endurance_regex.full_match(name)) {
 bool isRemaining = remaining_regex.full_match(name); // Check if it's a 'remaining life' type
 jglb["endurance_used"]["current_percent"] = isRemaining
 ? (normval <= 100 ? 100 - normval : 0) // Invert value if needed
 : normval; // Else use as-is
 return; // Exit after handling this attribute
}

// --- Temperature Attribute Parsing ---
// Matches all temperature-related labels like "Temperature_Celsius", "Temp_Internal", etc.
static const regular_expression temperature_regex(".*[Tt]emperature.*.*[Tt]emp.*");

// If attribute name matches temperature regex, store it under "temperature_celsius"
if (temperature_regex.full_match(name)) {

```

    jglb["temperature_celsius"] = normval; // Direct assignment as it's already normalized
    return;
}

// --- LBA Written / NAND Write Attribute Parsing ---
// Matches labels like "Host_Writes", "LBAs_Written", "Total_Writes_GiB", etc.
static const regular_expression lba_written_regex(

".*LBAs.*Written.*|.Host.*Writes.*|.Writes.*GiB.*|.Total.*Writes.*|.NAND.*Writes.*|.Program_
Page_Count.*"
);

// If matched, store raw value under "lbas_written"
if (lba_written_regex.full_match(name)) {
    jglb["lbas_written"] = normval; // Unit conversion (e.g., to GiB) can be added later if needed
    return;
}

// To handle edge cases
#ifdef NDEBUB
std::cerr << "Unmatched SMART attribute: " << name << " (ID: " << id << ", normval: " <<
normval << ")" << std::endl;
#endif

```

Testing & Environment Feedback (Self testing On my System)

Build Environment

- **System:** macOS Monterey (Darwin 21.6.0)
- **Hardware:** Apple SSD SM0128G (Model Family: Apple SD/SM/TS...E/F/G SSDs)
- **Toolchain:** Xcode CLI tools with `g++` (Apple clang)

Issue Encountered During Local Testing

Although the regex logic and attribute classification were implemented successfully, I faced challenges during local testing on my MacBook due to **compilation errors and runtime issues**:

Build Failure Summary

After modifying `ataprint.cpp`, `make` failed with:

```
error: 'continue' statement not in loop statement
error: function definition is not allowed here
```

- These were caused by misplaced `continue` statements outside any valid loop and improperly nested function definitions—likely due to a copy-paste or brace scope issue.

Cause

- Misplaced logic or braces led to function blocks being placed within other function bodies.

Run-Time Behavior

After building (pre-error versions), I attempted:

```
./smartctl -j -a /dev/disk0 | jq '.endurance_used'
```

But got:

```
zsh: no such file or directory: ./smartctl
```

This indicates:

- Either the binary wasn't built due to compile errors
- Or `make install` wasn't run to link it to `src/smartctl`

Also, due to macOS's **limited SMART visibility for Apple NVMe SSDs**, even successful runs often **lack vendor endurance attributes** (like `SSD_Life_Left`, etc.) in JSON output.

Github Draft PR Request for phase 1:

Summary

This pull request enhances the interpretation and normalization of SSD endurance-related SMART attributes in `smartmontools`, with a specific focus on standardizing lifetime usage reporting. The aim is to improve consistency across different vendor implementations by expanding detection heuristics and unifying the way values are presented in the JSON output.

Motivation

Vendors report SSD endurance metrics under wildly inconsistent attribute names in `drivedb.h`, including variations of "life used", "life remaining", "media wearout", and others. These discrepancies make it difficult for users and tools (e.g., Prometheus exporters) to meaningfully compare drive health, especially across fleets.

This PR addresses:

- The incomplete regex in `ataprint.cpp` used to identify endurance-related attributes.
- The lack of normalization between `% used` and `% remaining` values.
- The need to create a consistent `"endurance_used"` metric in the JSON output.

Changes Made

1. **Regex Expansion:**

- The `endurance_regex` in `ataprint.cpp` was expanded to match a broader and more accurate list of attribute labels found in `drivedb.h`.
- A secondary `remaining_regex` was introduced to detect attributes that report **remaining** life, allowing us to normalize them by computing `100 - normval`.

2. **Logic Adjustment:**

- If an attribute matches the `endurance_regex` and also matches the `remaining_regex`, the `normval` is subtracted from 100 to convert it to `% used`.
- All results are consistently reported under the `jglb["endurance_used"]["current_percent"]` key for downstream parsing.

Sample Test Cases

Attribute Name	Normval	Expected JSON Output
SSD_Life_Left	80	"endurance_used.current_percent": 20
Percent_Lifetime_Used	65	"endurance_used.current_percent": 65
Temperature_Celsius	35	"temperature_celsius": 35
Host_Writes_GiB	12456	"lbas_written": 12456
Unknown_Attr	50	(No JSON key, logged in debug only)

3. **Code Snippet Example:**

```
static const regular_expression endurance_regex(  
    "SSD_Life_Left.*|Wear_Leveling.*|Media_Wearout_Indicator.*|DriveLife_Used%.  
    *|DriveLife_Remaining%.*|Drive_Life_Used.*|Drive_Life_Remaining.*|"  
  
    "Lifetime_.*|Percent_Life_Used.*|Percent_Life_Remaining.*|Remaining_Life.*|  
    End_of_Life.*"  
);  
  
static const regular_expression remaining_regex(  
    ".*Left.*|.*Remaining.*|.*Media_Wearout_Indicator.*|.*End_of_Life.*"  
);  
  
if (id >= 100 && endurance_regex.full_match(name)) {  
    if (remaining_regex.full_match(name)) {  
        jglb["endurance_used"]["current_percent"] = (normval <= 100 ?  
100 - normval : 0);  
    } else {  
        jglb["endurance_used"]["current_percent"] = normval;  
    }  
    return;  
}
```

Future Work

The following improvements and extensions are planned for **Phase 2**:

1. Edge-Case Attribute Handling:

- Expand regex coverage to include rarely used or vendor-specific attributes like:
 - `Reserve_Blk_Remaining`, `Spare_Blocks_Remaining`,
`Spares_Remaining_Perc`, `Min_Spares_Remain_Perc`
- Introduce fallback heuristics or scoring logic for ambiguous or compound attributes (e.g., `Life_Curve_Status`, `Lifetime_NAND_Prg_GiB`).

2. Dynamic Label Inference:

- Develop a dynamic matching approach (e.g., substring or fuzzy matching) instead of static regex, to support future vendor formats without frequent updates.
- Use metadata like ID ranges, raw data types, or vendor flags from `drivedb.h` to aid classification.

3. Test Matrix Expansion:

- Evaluate endurance parsing against a broader set of SSDs (e.g., Samsung, WD, Seagate, Kingston) using virtualized Linux environments.
- Validate attribute normalization logic across NVMe and SATA interfaces using `/dev/sdX` and `/dev/nvmeX`.

4. Configuration Options:

- Add user-configurable options to allow overriding default regex behavior or remapping attributes in smartctl output.

5. Metric Output Enhancements:

- Support optional fields like:
 - `max_endurance_gb`, `lifetime_gb_written`,
`endurance_used_percent`
- Explore unit normalization for LBA writes (e.g., GiB, TB).

6. Upstream PR & Review Integration:

- Push incremental patches to smartmontools GitHub repo.
- Address feedback from maintainers and integrate with broader SMART attribute standardization efforts.
-

Learnings & Gaps

Learnings So Far

During Phase 1, I gained hands-on understanding of several smartmontools components and SSD attribute reporting concepts:

- **Regex-Based Label Normalization:**
Learned to design and maintain robust regular expressions for identifying endurance metrics with vendor-specific variations.
- **SMART Attribute Semantics:**
Developed familiarity with commonly used SSD attributes such as `Wear_Leveling_Count`, `Media_Wearout_Indicator`, `SSD_Life_Left`, and how vendors report usage vs. remaining lifetime.
- **smartmontools Internals:**
 - Navigated the structure and logic in `ataprint.cpp`, where SMART attributes are parsed and reported.
 - Understood how `drivedb.h` stores device-specific attribute mappings, and how parsing relies on ID-label-value triplets.

- Implemented JSON normalization using internal C++ data structures like `jglb` and `normval`.
- **Cross-Platform Considerations:**
Gained insight into platform-specific differences (e.g., macOS limitations vs. Linux's direct `/dev/sdX` access) in SMART visibility and build toolchain handling.

Gaps to Address

To complete the implementation and improve robustness, I aim to explore the following areas in Phase 2:

- **SMART Raw Value Interpretation:**
Deeper understanding of how raw SMART values (e.g., hex-encoded, vendor-scaled) map to usable endurance metrics across different firmware.
- **Dynamic Attribute Classification:**
Explore mechanisms beyond regex for fuzzy, dynamic, or metadata-based attribute inference.
- **Integration with Prometheus/Monitoring Tools:**
Investigate how `smartctl -j` output is consumed by exporters and what standard field naming conventions are expected in observability pipelines.
- **Unit Handling and Scaling:**
Study unit normalization for attributes like `Lifetime_Writes_GiB` vs. `LBAs_Written`, and clarify assumptions for conversion to TBW or % used.
- **Linux Build/Test Automation:**
Build testing environments (e.g., Docker or VM with Ubuntu) to validate changes on devices that provide richer SMART access than macOS.

Final Draft Pull Request – Phase 1: SSD Endurance Normalization in smartmontools - Done!

Summary

This draft PR implements logic to detect and normalize SSD endurance-related SMART attributes in smartmontools. It introduces regex-based matching for vendor-specific attribute labels, classifies attributes as representing used or remaining endurance, and outputs a unified endurance metric in the JSON report.

Motivation

SSD vendors report endurance metrics under widely varying attribute names in `drivedb.h`, such as variations of "life used", "life remaining", "media wearout", and others. This inconsistency complicates meaningful comparison of drive health, especially when monitoring large fleets using tools like Prometheus exporters.

This PR addresses the following issues:

- Incomplete regex matching in `ataprint.cpp` for endurance-related attributes.
- Lack of normalization between attributes reporting % used and % remaining endurance.
- Absence of a consistent "endurance_used" metric in the JSON output.

Key Changes

* **Expanded endurance_regex and remaining_regex:**

Covers a broader range of vendor-specific endurance labels such as `SSD_Life_Left`, `Wear_Leveling_Count`, `Percent_Lifetime_Remain`, and others.

* **Normalization Logic:**

- * If the attribute indicates *remaining* life, the value is inverted: $100 - x$.
- * Otherwise, the raw normalized value is used directly.
- * The normalized value is stored under `jglb["endurance_used"]["current_percent"]` in the JSON output.

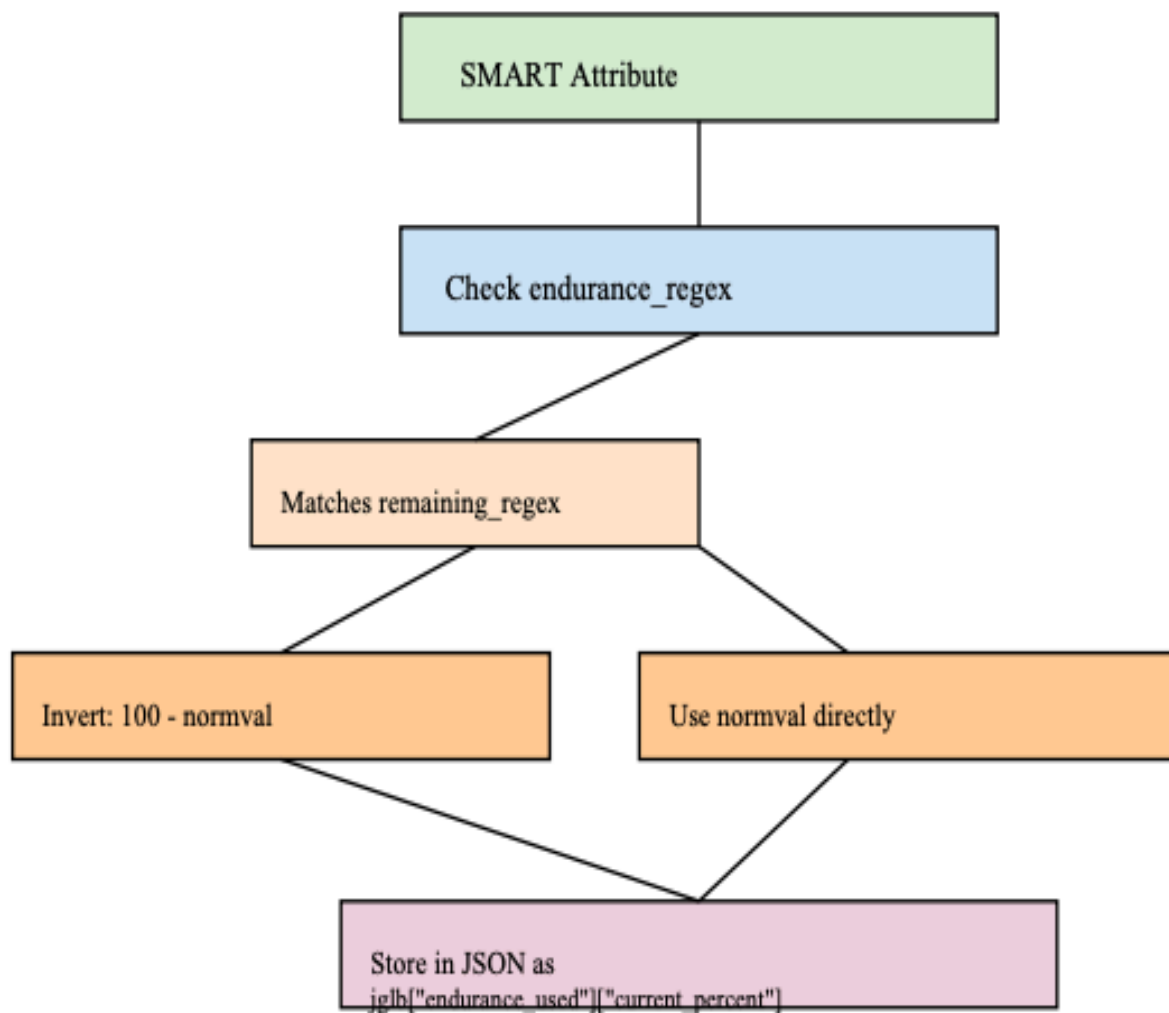
* **Improved Regex Coverage:**

Patterns mined from `drivedb.h` were added to capture less common labels like `Lifetime_Remaining%`, `End_of_Life`, etc.

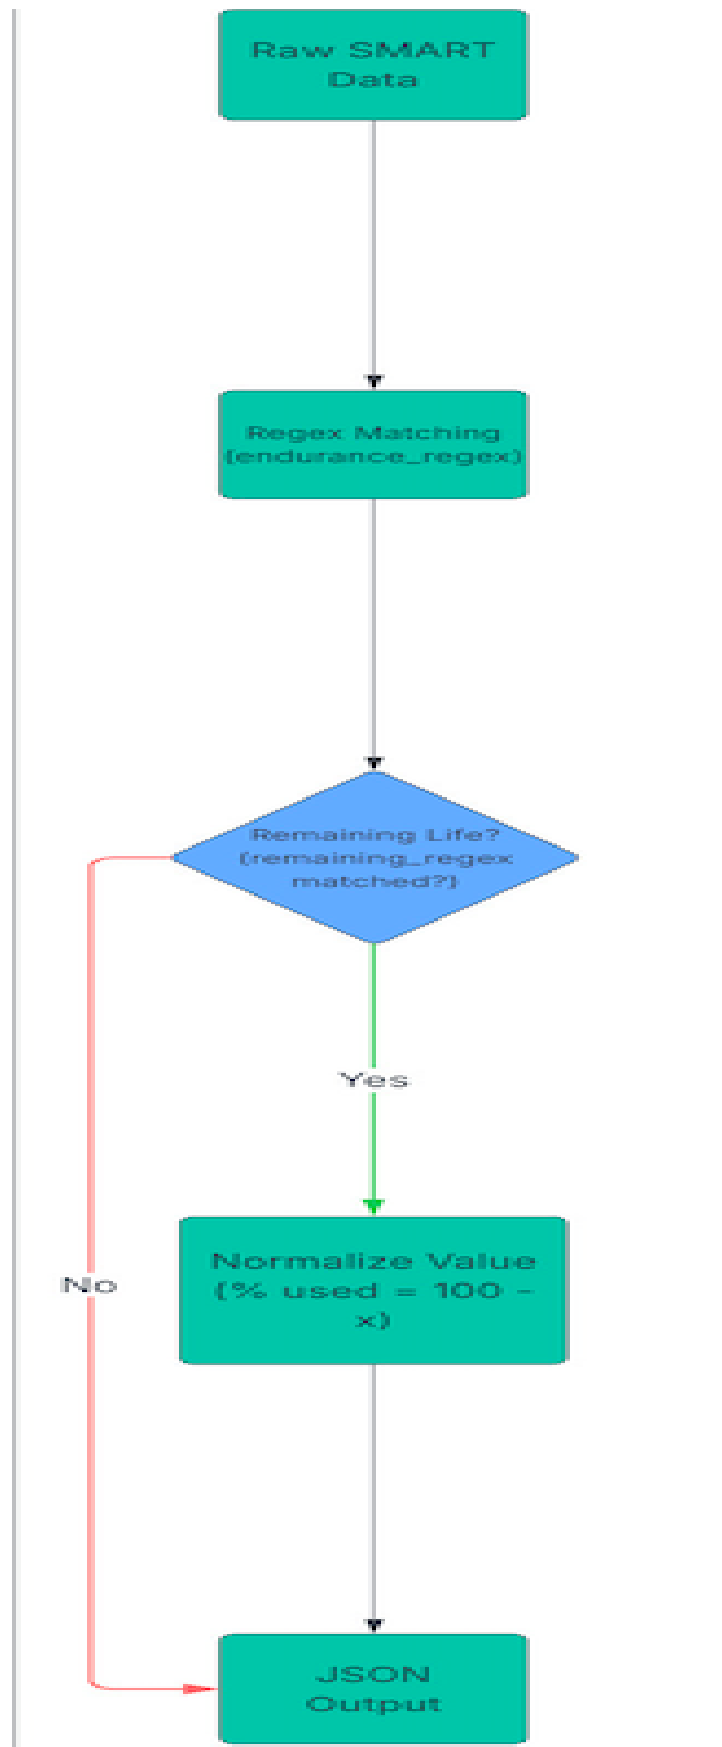
Execution Flow

The endurance normalization logic is summarized in the following flowcharts:

****Figure 1:**** Attribute matching and inversion logic.



****Figure 2:**** Overall parsing and output flow.



Changes in `src/ataprint.cpp`

For each SMART attribute `(id, name, normval)` processed, the following checks and actions are performed:

1. Endurance Check

- * If the attribute name matches `endurance_regex`:
 - * If it also matches `remaining_regex`, calculate endurance as:
`endurance_used.current_percent = 100 - normval`
 - * Else, use the normalized value directly:
`endurance_used.current_percent = normval`
 - * Exit the attribute check early after storing the endurance value.

2. Temperature Check

- * If the attribute name matches `temperature_regex`, store the value as:
`temperature_celsius = normval`

3. LBA Written Check

- * If the attribute name matches `lba_written_regex`, store the value as:
`lbas_written = normval`

4. Other Attributes

- * For attributes that do not match any of the above regex patterns:
 - * Log the unmatched attribute for debugging purposes (enabled only in debug builds).

****Success Path:**** JSON keys populated for matched attributes.

****Failure Path:**** Attributes unmatched — only logged in debug.

Sample Test Cases

Attribute Name	Normval	Expected JSON Output
SSD_Life_Left	80	"endurance_used.current_percent": 20
Percent_Lifetime_Used	65	"endurance_used.current_percent": 65
Temperature_Celsius	35	"temperature_celsius": 35
Host_Writes_GiB	12456	"lbas_written": 12456
Unknown_Attr	50	(No JSON key, logged in debug only)

Example Output

...

json

```
"endurance_used": {
  "current_percent": 76
}
```

...

Internals Touched

* ataprint.cpp:

- * Regex definitions for endurance and remaining life attributes.
- * Normalization logic implementation.
- * Conditional parsing extended for endurance, temperature, and LBAs written.

Notes

- This PR is in draft so regex coverage can be refined with more drive samples.
- macOS testing has limitations for Apple NVMe SSDs due to restricted SMART visibility.
- Future work: handle more edge-case labels (e.g., Reserve_Blk_Remaining).

To Do (Phase 2 / Final PR)

- * Add fuzzy matching to capture vendor-specific variations not covered by static regex.
- * Extend parsing logic to handle ambiguous attributes like Life_Curve_Status.
- * Improve testing coverage on Linux (Ubuntu) across SATA and NVMe SSDs.
- * Finalize PR by cleaning up commits and adding documentation comments.